

EE 5351 MP1

Lexi MacLean

1.

Given input matrices M , N , and P with dimensions $a \times b$, $b \times c$, and $a \times c$ respectively, each element in M will be read once for each column in P , and each element in N will be read once for each row in P . So, the total reads from each element in M is c , and the total reads from each element in N is a .

2.

A majority of the operations are done in the for loop, we are instructed to ignore the storage of the result (assumed in P , at least, the final storage is negligible). My for loop contains one statement:

```
c += M.elements[M.width * y + i] * N.elements[N.width * i + x];
```

This statement makes two reads from global memory: one from M and another from N . The statement also makes six floating point operations: four to compute the indices of the elements in M and N , one to multiply the two factors, and another to add them to the partial sum. The for loop itself also contributes an additional two floating point operations to increment and compare i . Ignoring operations out-side the for loop—which affect the ratio approaching 0 as the size of the input matrices increases—this makes for a 4 : 1 ratio of floating point operations to memory loads (or 3 : 1 if you ignore those made by the for loop).